

# Capstone Project

## A stock or Index Price Predictor

In [1]: *#Importing the major libraries*

```
import sklearn# Features various classification, regression and clustering algorithm
import numpy as np #Adding support for large, multi-dimensional arrays and matrices
import pandas as pd #Provides fast, flexible and expressive data structures
import matplotlib.pyplot as plt#A plotting library for Python and its numerical ecosystem
import datetime as dt #A module that supplies classes for manipulating dates and times
from matplotlib import style
import pandas_datareader.data as web #Getting data from the data provider's website
import seaborn as sns #A visualization library , provides high-level interface for making attractive plots
from math import sqrt #Will use this square root function to calculate RMSE
```

In [3]: *#Getting the data for a stock or Index*

```
start = dt.datetime(2000,1,1) #The start date for the data
end = dt.datetime(2017,8,31) # The end date for the purpose of this analysis
df= web.DataReader('GE', 'yahoo', start, end)# Getting the data for a stock in the form of a dataframe
print (df.head())# a representation of the first 5 set of data
print (df.tail())# a representation of the last 5 set of the data

df.to_csv('GE.csv') #converting the data retrieved into a csv file
```

	Open	High	Low	Close	Adj Close	Volume
Date						
2000-01-03	51.000000	51.229168	49.729168	50.000000	28.911114	22069800
2000-01-04	49.083332	49.333332	48.000000	48.000000	27.754667	22121400
2000-01-05	47.916668	49.000000	47.520832	47.916668	27.706484	27292800
2000-01-06	47.708332	48.979168	47.541668	48.557266	28.076889	19873200
2000-01-07	49.333332	50.625000	49.000000	50.437500	29.164083	20141400
	Open	High	Low	Close	Adj Close	Volume
Date						
2017-08-25	24.389999	24.600000	24.350000	24.490000	24.247725	22867800
2017-08-28	24.530001	24.670000	24.350000	24.469999	24.227922	23937600
2017-08-29	24.330000	24.459999	24.280001	24.440001	24.198219	23910100
2017-08-30	24.490000	24.490000	24.150000	24.280001	24.039803	33876000
2017-08-31	24.410000	24.700001	24.280001	24.549999	24.307131	55284300

In [4]: *#Visualizations*

```
#Get the adjusted close price for the stock for the stated period for analysis
Adj_close=web.DataReader('GE', 'yahoo', start, end)['Adj Close']

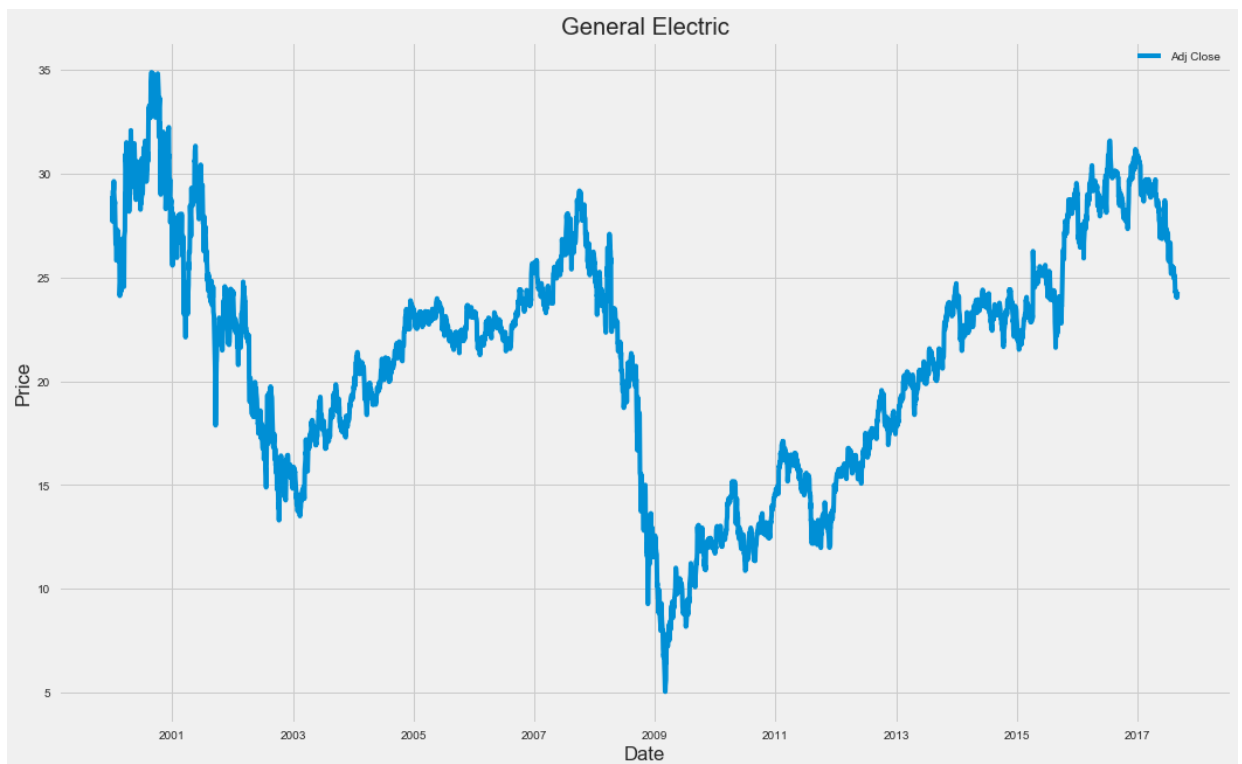
#print(Adj_close)

#graphing the Data

style.use('fivethirtyeight')

#Set the size of the graph display
plt.figure(figsize=(16,10))

plt.plot(Adj_close)
plt.legend()
plt.xlabel('Date')
plt.ylabel('Price')
lw=0.1
plt.title('General Electric')
plt.show()
```



## Generating Features

```

In [5]: def generate_features(df):
        """Generating features for a chosen stock or Index based on historical price
        the arguments passed to this function are df(dataframe with columns:"Open", "
        Function Returns:
            dataframe, data set with new features """
        df_new= pd.DataFrame()
        # The six original features

        df_new['open']= df['Open']
        df_new['Open_1']= df['Open'].shift(1) # This method shift index by 1, in orde

        df_new['close_1'] = df['Close'].shift(1)

        df_new['high_1']= df['High'].shift(1)

        df_new['low_1'] = df['Low'].shift(1)

        df_new['volume_1'] = df['Volume'].shift(1)

        # The 31 Features to Generate from the 6 original

        # Average price
        #The window sizes are rounded to 5 days, 21 days and 252 to represent the num

        df_new['avg_price_5']= pd.Series.rolling(df['Close'],window=5,center=False).m

        #rolling_mean calculates the moving average given a window {(Example [1,2,1,,

        df_new['avg_price_30']=pd.Series.rolling(df['Close'], window=21, center=False
        df_new['avg_price_365']=pd.Series.rolling(df['Close'], window=252, center=Fa

        #Ratio

        df_new['ratio_avg_price_5_30'] = df_new['avg_price_5'] / df_new['avg_price_30
        df_new['ratio_avg_price_5_365'] = df_new['avg_price_5'] / df_new['avg_price_3
        df_new['ratio_avg_price_30_365'] = df_new['avg_price_30'] / df_new['avg_price

        # average volume
        df_new['avg_volume_5'] =pd.Series.rolling(df['Volume'], window=5, center=False
        df_new['avg_volume_30'] = pd.Series.rolling(df['Volume'], window=21, center=F
        df_new['avg_volume_365'] =pd.Series.rolling(df['Volume'], window=252, center=
        df_new['ratio_avg_volume_5_30'] =df_new['avg_volume_5'] / df_new['avg_volume_
        df_new['ratio_avg_volume_5_365'] =df_new['avg_volume_5'] / df_new['avg_volum
        df_new['ratio_avg_volume_30_365'] =df_new['avg_volume_30'] / df_new['avg_volum

        # standard deviation of prices

        df_new['std_price_5'] = pd.Series.rolling(df['Close'], window=5, center=False

        # rolling_mean calculates the moving standard deviation given a window
        df_new['std_price_30'] =pd.Series.rolling(df['Close'], window=21, center=False
        df_new['std_price_365'] =pd.Series.rolling(df['Close'], window=252, center=Fa
        df_new['ratio_std_price_5_30'] =df_new['std_price_5'] / df_new['std_price_30'

```

```

df_new['ratio_std_price_5_365'] =df_new['std_price_5'] / df_new['std_price_365']
df_new['ratio_std_price_30_365'] =df_new['std_price_30'] / df_new['std_price_365']

# standard deviation of volumes
df_new['std_volume_5'] =pd.Series.rolling(df['Volume'], window=5,center=False)
df_new['std_volume_30'] = pd.Series.rolling(df['Volume'], window=21, center=False)
df_new['std_volume_365'] =pd.Series.rolling(df['Volume'], window=252, center=False)
df_new['ratio_std_volume_5_30'] =df_new['std_volume_5'] / df_new['std_volume_30']
df_new['ratio_std_volume_5_365'] =df_new['std_volume_5'] / df_new['std_volume_365']
df_new['ratio_std_volume_30_365'] =df_new['std_volume_30'] / df_new['std_volume_365']

# return
df_new['return_1'] = ((df['Close'] - df['Close'].shift(1))/ df['Close'].shift(1))
df_new['return_5'] = ((df['Close'] - df['Close'].shift(5))/ df['Close'].shift(5))
df_new['return_30'] = ((df['Close'] -df['Close'].shift(21)) / df['Close'].shift(21))
df_new['return_365'] = ((df['Close'] - df['Close'].shift(252)) / df['Close'].shift(252))
df_new['moving_avg_5'] =pd.Series.rolling(df_new['return_1'], window=5, center=False)
df_new['moving_avg_30'] = pd.Series.rolling(df_new['return_1'], window=21, center=False)
df_new['moving_avg_365'] = pd.Series.rolling(df_new['return_1'], window=252, center=False)

# the target
df_new['close'] = df['Close']
df_new = df_new.dropna(axis=0) # This will drop rows with any N/A value, which is the target

return df_new

```

## Applying feature engineering strategy to the GE data

```

In [6]: raw_data=df
data=generate_features(raw_data)

data.round(decimals=2).head(3)

```

```

Out[6]:

```

	open	Open_1	close_1	high_1	low_1	volume_1	avg_price_5	avg_price_30	avg_price_365
<b>Date</b>									
<b>2001-01-03</b>	44.25	46.75	43.75	46.88	42.62	36837700.0	47.52	50.70	51
<b>2001-01-04</b>	47.31	44.25	47.81	47.94	43.81	39205800.0	47.22	50.55	51
<b>2001-01-05</b>	47.75	47.31	48.06	48.75	47.12	26926400.0	47.20	50.38	51

3 rows × 38 columns

## Selecting data for training and testing

```
In [7]: import datetime

#All the fields in the dataframe 'data' are feature columns while 'Close' is the
start_train =datetime.datetime(2000,1,1,0,0)
end_train =datetime.datetime(2017,8,31,0,0)
data_train =data.loc[start_train:end_train]

#print (data_train)
```

```
In [8]: #Partitioning the data between training and testing sets

# We will be using the TimeSeriesSplit

"""TimeSeriesSplit Provides train/test indices to split time series data samples
in train/test sets. In each split, test indices must be higher than before, and
validator is inappropriate."""

from sklearn.model_selection import TimeSeriesSplit

X,y= data_train.iloc[:, 1:].values, data_train.iloc[:,0].values #Assigning values
#print(X)
#print(y)

dataList=TimeSeriesSplit(n_splits=3)

print(dataList)

for train_index, test_index in dataList.split(X):
    print("TRAIN:", train_index, "TEST:", test_index)
    X_train, X_test = X[train_index], X[test_index]
    y_train, y_test = y[train_index], y[test_index]

X_train.shape
y_train.shape
y_test.shape
```

```
TimeSeriesSplit(n_splits=3)
TRAIN: [  0   1   2 ..., 1045 1046 1047] TEST: [1048 1049 1050 ..., 2093 209
4 2095]
TRAIN: [  0   1   2 ..., 2093 2094 2095] TEST: [2096 2097 2098 ..., 3141 314
2 3143]
TRAIN: [  0   1   2 ..., 3141 3142 3143] TEST: [3144 3145 3146 ..., 4189 419
0 4191]
```

```
Out[8]: (1048,)
```

## Preprocessing the data

```
In [9]: from sklearn.preprocessing import StandardScaler

sc=StandardScaler() #Using stantardization

sc.fit(X_train)

#Rescaling both sets using the trained scaler; feature scaling is a crucial steep

X_scaled_train=sc.transform(X_train)
X_scaled_test=sc.transform(X_test)

X_test.shape
```

Out[9]: (1048, 37)

## Implementation of the Linear Regression Model

In [10]: *#Using a Stochastic Gradient Descent(SGD)-based Linear Rregression*

```
from sklearn import linear_model
from sklearn.grid_search import GridSearchCV
from sklearn.metrics import mean_squared_error , r2_score

#Setting the optimal set of parameters for the regression
Lr=linear_model.SGDRegressor(alpha=0.0001, penalty='l2', n_iter=1000, eta0=0.01)

"""Implementing the Optimization parameter 'grid search' that can further improve
the Optimal combinaton of hyperparameters"""

#A list of dictionaies that have the parameters that we want to tune
param_grid={"alpha":[3e-06, 1e-05, 3e-5],"eta0":[0.01,0.03,0.1]}

grid_search=GridSearchCV(Lr,param_grid=param_grid, cv=5, scoring='neg_mean_square
grid_search.fit(X_scaled_train, y_train)

print(grid_search.best_params_)
```

C:\Users\delis\Anaconda3\lib\site-packages\sklearn\cross\_validation.py:44: DeprecationWarning: This module was deprecated in version 0.18 in favor of the model\_selection module into which all the refactored classes and functions are moved. Also note that the interface of the new CV iterators are different from that of this module. This module will be removed in 0.20.

"This module will be removed in 0.20.", DeprecationWarning)

C:\Users\delis\Anaconda3\lib\site-packages\sklearn\grid\_search.py:43: DeprecationWarning: This module was deprecated in version 0.18 in favor of the model\_selection module into which all the refactored classes and functions are moved. This module will be removed in 0.20.

DeprecationWarning)

```
{'alpha': 3e-06, 'eta0': 0.03}
```

### ***Implementing the Prediction and Measuring its performance***

```

In [11]: Lr_best = grid_search.best_estimator_

y_train_pred=Lr_best.predict(X_scaled_train)
y_test_pred1=Lr_best.predict(X_scaled_test)

#rmse=sqrt(mean_squared_error(y_test,y_test_pred1))

#The Mean Squared Error and R^2

#Coefficient of Determination (R^2, this is the fraction of response variance tha
#
print('MSE train: %.3f, test:%.3f' %(mean_squared_error(y_train, y_train_pred), m
print('RMSE train: %.3f, test:%.3f' %(sqrt(mean_squared_error(y_train, y_train_pr
print('R^2 train: %.7f, test: %.7f' % (r2_score(y_train, y_train_pred), r2_score(
print(y_test_pred1)

```

MSE train: 0.068, test:0.027

RMSE train: 0.260, test:0.166

R^2 train: 0.9992260, test: 0.9951264

[ 23.27383518 23.4668783 23.65255756 ..., 24.4542678 24.36308018  
24.34792929]



In [11]: *#Visualizations of the Residuals of the Prediction*


```
plt.figure(figsize=(12,9))

plt.scatter(y_train_pred, y_train_pred - y_train, c='black', marker='o', s=35, alp

plt.scatter(y_test_pred1, y_test_pred1 - y_test, c='lightgreen', marker='s', s=35,

plt.title('Residuals of the Prediction')
plt.xlabel('Predicted values')
plt.ylabel('Residuals')
plt.legend(loc='upper left')
plt.hlines(y=0, xmin=-10, xmax=100, lw=2, color='red')
plt.xlim([-10, 50])

plt.show()
```



## Implementation of the Random Forest Algorithm

```
In [12]: from sklearn.ensemble import RandomForestRegressor

        """A random forest, is an ensemble of multiple decision tree. Here we are subdividing
           will become more manageable."""

        #Assigning parameters for tuning the regressor

        forest=RandomForestRegressor(n_estimators=1000, random_state=1, n_jobs=1, max_depth=5)

        forest.fit(X_scaled_train, y_train)

        y_train_pred=forest.predict(X_scaled_train)
        y_test_pred2=forest.predict(X_scaled_test)

        #The Mean Squared Error and R^2
        print('MSE train: %.3f, test: %.3f' % (mean_squared_error(y_train, y_train_pred), mean_squared_error(y_test, y_test_pred2)))
        print('RMSE train: %.3f, test: %.3f' % (sqrt(mean_squared_error(y_train, y_train_pred)), sqrt(mean_squared_error(y_test, y_test_pred2))))
        print('R^2 train: %.3f, test: %.3f' % (r2_score(y_train, y_train_pred), r2_score(y_test, y_test_pred2)))
```

```
MSE train: 0.012, test:0.050
RMSE train: 0.112, test:0.224
R^2 train: 1.000, test: 0.991
```

In [13]: *#Visualizations of the Residuals of the Prediction*


```
plt.figure(figsize=(14,9))

plt.scatter(y_train_pred, y_train_pred - y_train, c='black', marker='o', s=35, alp

plt.scatter(y_test_pred2, y_test_pred2 - y_test, c='lightgreen', marker='s', s=35,

plt.title('Residuals of the Prediction')
plt.xlabel('Predicted values')
plt.ylabel('Residuals')
plt.legend(loc='upper left')
plt.hlines(y=0, xmin=-10, xmax=50, lw=2, color='red')
plt.xlim([-10, 50])

plt.show()
```



## Implementation of the Support Vector Machine(SVM) Algorithm

```
In [14]: from sklearn.svm import SVR

svr=SVR(kernel='linear', gamma=100.0, C=1.0)
svr.fit(X_scaled_train, y_train)

y_train_pred=svr.predict(X_scaled_train)
y_test_pred3=svr.predict(X_scaled_test)

#The Mean Squared Error, RMSE and R^2
print('MSE train: %.3f, test:%.3f' %(mean_squared_error(y_train, y_train_pred), m
print('RMSE train: %.3f, test:%.3f' %(sqrt(mean_squared_error(y_train, y_train_pr
print('R^2 train: %.7f, test: %.7f' % (r2_score(y_train, y_train_pred), r2_score(

#print(y_test_pred2)

MSE train: 0.067, test:0.025
RMSE train: 0.260, test:0.157
R^2 train: 0.9992302, test: 0.9956108
```

In [15]: *#Visualizations of the Residuals of the Prediction*


```
plt.figure(figsize=(12,9))

plt.scatter(y_train_pred, y_train_pred - y_train, c='black', marker='o', s=35, alp

plt.scatter(y_test_pred3, y_test_pred3 - y_test, c='lightgreen', marker='s', s=35,

plt.title('Residuals of the Prediction')
plt.xlabel('Predicted values')
plt.ylabel('Residuals')
plt.legend(loc='upper left')
plt.hlines(y=0, xmin=-10, xmax=100, lw=2, color='red')
plt.xlim([-10, 50])

plt.show()
```



## The Benchmark (Using a Naive Forecast)

```

In [16]: """I will be using a Naive forecast, this forecasting method is suitable for Time
we will be using a seasonal naive approach with the autoregressive integrated mov
generalization of an autoregressive moving average (ARMA) model. Both of these mo
better understand the data or to predict future points in the series (forecasting

from statsmodels.tsa.arima_model import ARIMA

series=Adj_close #Using the adjusted close for the dates(January 1, 2000 - August
X=series.values #Getthng the Adjusted closing prices ffrom the data
#print(X)

#Partitioning the data for trainign and testing
size=int(len(X)*0.77)

train, test=X, X[size:len(X)]

history=[values for values in train] #parsing through the training data

predict=list()

for t in range(len(test)):
    model=ARIMA(history, order=(5,1,0))#Initializing the ARIMA module; 5,1,0 repr
    model_fit=model.fit(dis=0)
    output=model_fit.forecast()#Using the ARIMA for making forecasts
    y_t=output[0]
    predict.append(y_t)
    obs = test[t]
    history.append(obs)

error=mean_squared_error(test, predict)
RMSE=sqrt(mean_squared_error(test,predict))
R_2 = r2_score(test, predict)

print('Test MSE: %.3f' %error)
print('Test RMSE: %.3f'%RMSE)
print('Test R_2: %.3f' %R_2)

```

C:\Users\delis\Anaconda3\lib\site-packages\statsmodels\compat\pandas.py:56: FutureWarning: The pandas.core.datetools module is deprecated and will be removed in a future version. Please use the pandas.tseries module instead.

from pandas.core import datetools

C:\Users\delis\Anaconda3\lib\site-packages\statsmodels\base\model.py:496: ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check mle\_retvals

"Check mle\_retvals", ConvergenceWarning)

C:\Users\delis\Anaconda3\lib\site-packages\statsmodels\base\model.py:496: ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check mle\_retvals

"Check mle\_retvals", ConvergenceWarning)

Test MSE: 0.089

Test RMSE: 0.299

Test R\_2: 0.990

## Plot Comparison

```
In [17]: import seaborn as sns

plt.figure(figsize=(16,9))
sns.set_context(font_scale=3)

plt.xlim([0.0, 1100])
plt.ylim([0.0, 60])
lw=1
plt.title('Stock price prediction Models vs Benchmark[Naive-Forecast(ARIMA)]')

# Set figure width to 12 and height to 9

plt.xlabel('Number of Testing instance')
plt.ylabel('Price')

plt.plot(y_test_pred1, c='black',marker='o', lw=0.75,alpha=0.5, label='Linear Reg
plt.plot(y_test_pred2, c='red', marker='x', lw=3, label='RandomForest')
plt.plot(y_test_pred3, c='yellow', lw=1, label='SVR')
plt.plot(predict,c='green', lw=0.5, label='Benchmark-Naive-Forecast(ARIMA)')

plt.legend(loc='lower left')
plt.show()
```

